

Dissecting Netflix’s Self-Preferencing: Evidence from Viewer-Level Data^{*}

Tin Cheuk Leung[†]

Wake Forest University

Shi Qi[‡]

William & Mary

Koleman Strumpf[§]

Wake Forest University

June 2026

Abstract

This paper studies self-preferencing in subscription video streaming. We assemble a new dataset linking Netflix’s U.S. catalog, Top 10 rankings, Wikipedia page views, and device-level streaming records. Conditional on external popularity and other observables, Netflix Originals are significantly more likely to appear in the Top 10 than comparable non-original titles, especially among series. Using a matched staggered difference-in-differences design around first Top 10 entry, we show that Top 10 placement causally increases subsequent viewership. The effect is weaker among devices with extensive prior Netflix viewing, consistent with personalized recommendations substituting for generic rankings. These findings provide evidence of platform gatekeeping in digital media.

JEL Classification: D22, K21, L40, L82, M21

Keywords: Self-Preferencing, Netflix, Digital Platforms, Platform Bias

^{*}We are grateful to seminar participants at Bates College, CREED, the Mallen Conference, and SEA for helpful comments. We thank Andy Shen for his research assistance. We thank the NET Institute www.NETinst.org and the discretionary funding from Wake Forest for financial support. This research includes content supplied by Samba TV; Copyright ©Samba TV, 2026. All rights reserved. The usual disclaimer applies.

[†]Contact: Department of Economics, Kirby Hall, Wake Forest University, Box 7505, Winston-Salem, NC 27109. Email: leungtc@wfu.edu

[‡]Email: sqi01@wm.edu

[§]Email: strumpks@wfu.edu

1 Introduction

Digital platforms increasingly determine which products, services, and content consumers see first. Search rankings, recommendation rows, app-store placement, marketplace buy boxes, and curated popularity lists are not merely passive ordering devices. They allocate attention in markets where attention is scarce and where small changes in visibility can redirect demand. This control over visibility has made self-preferencing a central issue in competition policy. When a platform both operates an interface and supplies its own products or content on that interface, it may have the incentive and ability to steer users toward affiliated offerings rather than observationally similar third-party alternatives. The competitive significance of such steering depends on market power. Preferential placement by a small platform may be commercially relevant, but it has limited competitive significance if users and suppliers can easily bypass the platform. By contrast, when consumers rely on a platform for discovery and content providers rely on it for access to demand, rankings and recommendations can become instruments of gatekeeping. Whether Netflix has meaningful market power is therefore central to why self-preferencing in streaming is an economically relevant object of study.

Recent policy debates reflect this concern. In the United States, the American Innovation and Choice Online Act (AICOA) would prohibit covered platforms from unfairly preferencing their own products, services, or lines of business and from treating their own offerings more favorably in user interfaces, including search or ranking functionality.¹ In Europe, the Digital Markets Act (DMA) imposes direct obligations on designated gatekeepers, including a prohibition on treating the gatekeeper’s own services and products more favorably in ranking than similar third-party services and products.² The debate has continued in state-level proposals such as California’s Blocking Anticompetitive Self-preferencing by Entrenched Dominant platforms, or BASED Act. The recent discussion around its committee defeat framed self-preferencing as a live dispute over whether platform-neutrality rules protect competition or constrain product design.³ These debates

¹See S.2992, American Innovation and Choice Online Act, 117th Congress, available through Congress.gov.

²See Regulation (EU) 2022/1925, Article 6(5), available through EUR-Lex.

³See ITIF press release, April 21, 2026.

have focused primarily on search engines, app stores, e-commerce marketplaces, operating systems, and online advertising. Yet the same visibility problem arises in media markets, where platforms decide how content is ranked, recommended, and displayed.

Video streaming is an important but understudied setting for these issues. Streaming platforms do not simply host content. They increasingly act as attention intermediaries, organizing large catalogs through homepages, personalized recommendation rows, search results, autoplay previews, genre shelves, curated lists, and non-personalized rankings such as Netflix’s Top 10 list. Despite this central role in discovery, empirical evidence on preferential treatment in streaming remains limited relative to search engines, app stores, and e-commerce platforms. This gap matters because streaming demand is highly visibility-sensitive. Consumers face large catalogs, low marginal prices, and high search costs, so interface placement can materially affect what is watched. The streaming context also differs from transaction-based marketplaces. Netflix earns subscription revenue rather than per-title commissions, and it competes with outside studios while also producing its own originals. This structure creates a distinct incentive to allocate attention toward proprietary content that is stable, exclusive, and useful for subscriber retention.

Netflix is also an economically meaningful setting because market participants increasingly debate whether it has gatekeeper power in streaming. The recent Netflix-Warner Bros. Discovery transaction debate illustrates this perceived market power. Netflix’s proposed acquisition reportedly valued Warner Bros. Discovery at approximately \$82.7 billion and generated commentary that a combined firm could become the dominant global streaming platform, with analysts estimating a potential 30–40% U.S. streaming-market share and substantial antitrust scrutiny.⁴ Although the deal ultimately failed, the public debate is informative. Concerns centered on scale, subscriber reach, bargaining leverage, data advantages, and control of premium content libraries. These are precisely the conditions under which preferential visibility can have stronger competitive consequences. If Netflix is viewed by some observers as a gatekeeper for streaming demand, then understanding whether it advantages its own content becomes directly relevant to industrial

⁴See Corey Martin, “Netflix-Warner Bros. Deal: Streaming Superpower Or Streaming Monopoly?,” *Forbes*, December 5, 2025. Netflix later declined to match Paramount Skydance’s higher bid, ending its pursuit of Warner Bros. Discovery. See Megan Cerullo, “Netflix drops \$83 billion bid for Warner Bros. Discovery, paving way for Paramount Skydance deal,” CBS News, February 27, 2026.

organization and antitrust policy.

This paper studies self-preferencing and ranking effects on Netflix. We ask three questions. First, conditional on observable popularity and title characteristics, are Netflix Originals more likely than non-original titles to receive prominent Top 10 visibility? Second, does Top 10 inclusion causally increase subsequent demand, or does it simply summarize titles that would have become popular anyway? Third, do users respond similarly to these rankings, or does responsiveness vary with their prior Netflix histories?

Answering these questions requires data that are rarely available for streaming platforms. We build a weekly panel of the Netflix U.S. catalog from 2016 through early 2025, tracking title entry, exit, original status, availability, and other characteristics. Constructing this catalog is itself non-trivial. Streaming libraries are dynamic, titles enter and exit frequently, original and licensed content must be distinguished, and public archival coverage is incomplete. We combine this catalog with Netflix’s official Top 10 rankings, title-level Wikipedia page views, and Samba TV’s device-level streaming data from U.S. smart TVs. The two demand measures are complementary. Wikipedia page views provide a broad, externally observable measure of title-level attention and discovery. Device-level streaming records provide direct behavioral evidence on viewing, allowing us to study granular responses to rankings and heterogeneity across devices.

We find evidence consistent with Netflix self-preferencing. Conditional on Wikipedia page views and other observable controls, Netflix Originals are more likely to appear in the Top 10 than comparable non-original titles. The pattern is especially pronounced for series, which are also the content type most closely linked to subscriber retention in our mechanism analysis. We also find that Netflix appears to deprioritize titles that are close to leaving the platform, consistent with incentives to avoid directing attention toward content with limited remaining licensing windows.

Rankings also matter for demand. Using a matched staggered difference-in-differences design around first Top 10 entry, we find that inclusion in the Top 10 generates large increases in subsequent viewership. The effects appear in both Wikipedia traffic and streaming-device outcomes, and they are larger for series and for Netflix Originals. These results indicate that the Top 10 list is not merely a descriptive measure of contempora-

neous popularity. It is also a visibility mechanism that can amplify demand after a title receives prominent placement.

The device-level evidence reveals substantial heterogeneity in these ranking effects. In a 10 percent random sample of devices with at least eight active Netflix weeks, we estimate device-specific responsiveness to lagged Top 10 status. The average effect remains positive, but many devices have coefficients close to zero, indicating little measurable response to the generic ranking. Moreover, responsiveness is negatively correlated with prior viewing history, including total viewing duration and the number of active Netflix weeks. This pattern is descriptive, but it is consistent with a signal-substitution interpretation: for new or infrequent viewers, a common Top 10 list may provide useful information about what to watch, while experienced viewers may rely more on personalized recommendations, accumulated viewing histories, and prior knowledge of the catalog. In this sense, the heterogeneity results point indirectly to the growing economic importance of algorithmic personalization.

The remainder of the paper proceeds as follows. Section 2 reviews the related literature. Section 3 describes Netflix’s institutional background, catalog strategy, and content discovery interface. Section 4 describes the data and validates Wikipedia page views as a proxy for title-level interest. Section 5 tests for self-preferencing in Top 10 placement. Section 6 examines the subscriber-retention mechanism using content exits and device-level churn. Section 7 estimates the causal impact of Top 10 visibility and studies heterogeneity in device-level responsiveness. Section 8 discusses policy implications, and Section 9 concludes.

2 Literature Review

2.1 Self-Preferencing and Platform Gatekeeping

This paper contributes to the growing literature on self-preferencing and platform gatekeeping. Much of this work studies hybrid platforms that both operate a marketplace and compete with sellers or content providers on that marketplace. In e-commerce and related settings, researchers have examined whether platforms steer attention toward their

own products or favored partners through search rankings, recommendations, and default placements (Farronato et al., 2023; Reimers and Waldfogel, 2023; Jürgensmeier et al., 2024). Recent evidence from Amazon shows own-product ranking advantages that declined after its DMA gatekeeper designation (Waldfogel, 2024). Other work studies the welfare and competitive implications of platform entry, private-label competition, and self-preferencing (Gutiérrez, 2022; Lam, 2021; Lee and Musolff, 2023; Hagiu et al., 2022). Beyond retail platforms, related studies examine search and ranking bias in app stores, travel platforms, and music streaming (Teng, 2023; Hunold et al., 2020; Cure et al., 2022; Aguiar et al., 2021).

The streaming setting differs from these canonical cases. Netflix does not earn a separate transaction margin from each view, and users face no per-title price at the point of consumption. The platform’s incentive to favor its own titles instead operates through attention allocation, subscriber retention, and the management of dependence on licensed content. This distinction makes streaming a useful setting for studying how self-preferencing works when the relevant margin is engagement rather than a transaction.

2.2 Rankings, Recommendations, and Popularity Signals

The paper also relates to research on rankings, recommendation systems, and information design. Prominence can affect consumer consideration and demand even when products are otherwise available (Armstrong and Zhou, 2011; Hagiu and Jullien, 2014; de Cornière and Taylor, 2019). In digital media, platform-operated recommendations and playlists can substantially affect consumption outcomes (Aguiar and Waldfogel, 2021; Aguiar et al., 2021). Theoretical work on streaming recommendations further shows that platforms may have incentives to bias recommendations when they control both the ranking technology and the content portfolio (Bourreau and Gaudin, 2022).

Our setting adds evidence on a simple and widely visible form of information design: a public Top 10 ranking. Such rankings may convey information about popularity, reduce search costs, and create social-learning or salience effects. The empirical challenge is that popular titles are more likely to be ranked even absent any causal effect of the ranking. By combining pre-treatment matching with staggered event-study variation around first

Top 10 entry, we separate the demand effect of visibility from pre-existing popularity more directly than raw ranking correlations would allow.

2.3 Media, Streaming, and Digital Content Markets

The paper also contributes to the economics of digital media and streaming. Prior work has studied digital music consumption, playlist promotion, global Netflix availability, and consumer demand for streaming services (Aguilar and Martens, 2016; Aguilar and Waldfogel, 2018, 2021; McKenzie et al., 2019). Other work examines how streaming availability affects adjacent markets such as film sales and piracy, and how exclusive contracts shape competition among video platforms (Danaher and Doosti, 2025; Frick et al., 2023; Yuan, 2025). These studies show that streaming platforms matter for content distribution and consumption, but direct evidence on platform-controlled visibility within video streaming remains limited.

Our empirical setting is useful because Netflix combines a large catalog, a visible non-personalized Top 10 ranking, and a growing portfolio of proprietary originals. This combination allows us to study both sides of platform gatekeeping: whether the platform allocates visibility differentially across affiliated and unaffiliated content, and whether that visibility changes subsequent demand.

2.4 Our Contribution

The paper makes four contributions. First, it studies self-preferencing in subscription video streaming rather than in classic retail, search, or app-store settings. Second, it combines title-level attention data from Wikipedia with device-level behavioral data, allowing us to measure both broad public interest and actual viewing responses. Third, it connects preferential placement to downstream demand effects by showing that Top 10 rankings are causal demand shifters rather than merely descriptive popularity signals. Fourth, it documents substantial heterogeneity in ranking responsiveness and links that heterogeneity to prior watching history. This evidence points to an interaction between public popularity signals and personalized recommendation systems, a margin that is central to modern digital platforms but difficult to observe directly.

3 Institutional Background

Netflix is the leading firm in the subscription-based video streaming market, with more than 260 million global subscribers as of early 2025. Netflix operates under a flat-rate subscription model, where users face zero marginal cost for consuming additional content. As a result, the platform’s primary objective is not to maximize revenue per transaction, but rather to attract and retain subscribers over time. This business model places a premium on sustained user engagement and reduces the direct link between individual viewing decisions and platform revenue.

A central feature of Netflix’s platform is its content discovery system, which governs how users are exposed to available titles. Recommendations take multiple forms. Personalized rows, such as “Because You Watched” or “Recommended for You,” are tailored to individual viewing histories and are dynamically updated using proprietary algorithms. In addition, Netflix prominently displays a “Continue Watching” row, which encourages users to resume previously viewed content. These personalized features rely heavily on accumulated user data and are designed to match content to individual preferences.

Alongside these personalized recommendations, Netflix also employs non-personalized forms of content promotion. In 2021, the platform introduced a weekly “Top 10” ranking, which highlights the most-watched movies and television series within each country. Unlike personalized recommendations, the Top 10 list is identical across users within a given region and is prominently displayed on the home interface. As such, it serves as a common visibility mechanism that can influence viewing decisions, particularly for users with limited viewing histories or weaker preference signals.

Over the past decade, Netflix has increasingly shifted toward producing and distributing its own original content. Beginning with early successes such as *House of Cards*, the platform has expanded its portfolio of originals to include high-profile series such as *Stranger Things*, *The Crown*, and *Bridgerton*, as well as a growing slate of original films. This transition reflects a strategic move away from reliance on licensed content, which is subject to expiration and competitive withdrawal as rival platforms develop their own streaming ecosystems.

This distinction between original and licensed content has important implications for

platform incentives. Licensed titles are typically governed by fixed-term contracts and may exit the platform when agreements expire. Such exits can reduce the attractiveness of the catalog and potentially lead to subscriber churn, particularly for popular serialized content. By contrast, Netflix Originals are fully controlled by the platform, are globally licensed, and can be promoted without contractual constraints. As a result, Originals represent a more stable and strategically valuable asset for sustaining long-term engagement.

Importantly, Netflix retains full control over the algorithms and interface elements that determine content visibility, including both personalized recommendations and curated rankings such as the Top 10 list. This control enables the platform to influence the allocation of user attention across titles. While such influence may reflect underlying consumer demand, it also creates the possibility that Netflix systematically favors its own content in ways that are not fully explained by observable measures of popularity.

These features of the Netflix platform, its subscription-based revenue model, reliance on algorithmic content discovery, and increasing vertical integration into content production, create a setting in which self-preferencing may arise through the strategic allocation of visibility rather than through price or transaction-level distortions. This institutional context motivates our empirical analysis of whether Netflix differentially promotes its own content and the extent to which such promotion affects viewer behavior.

4 Data

Our analysis combines four data components: a weekly panel of Netflix’s U.S. catalog, Netflix’s official Top 10 rankings, title-level Wikipedia page views, and device-level streaming records from Samba TV. The catalog and ranking data describe the titles Netflix makes available and promotes. Wikipedia page views provide an external measure of title-level attention. The device-level data measure realized viewing behavior and allow us to validate the attention proxy and study heterogeneous responses to ranking visibility.

4.1 Netflix’s Weekly Catalog

We construct a weekly panel of the Netflix U.S. catalog from 2016 through early 2025. The panel tracks title availability, entry and exit dates, Netflix Original status, genre, country of origin, primary language, and other metadata used in the empirical specifications. Catalog construction is non-trivial because streaming availability changes frequently, public archives are incomplete, and title metadata must be harmonized across sources.

We identify entry dates using daily records from New on Netflix (<https://usa.newonnetflix.info/>), which reports newly added titles in the U.S. catalog. We identify exit dates and title metadata using the uNoGS API (<https://unogs.com/>). Starting from a full catalog snapshot in March 2025, we reconstruct historical availability by moving backward week by week. In each week, titles that had not yet entered are excluded, while titles with expiration dates after the current week remain in the catalog. This procedure yields a consistent title-week panel of Netflix U.S. availability.

For each title-week, we record whether the title is a Netflix Original, whether it is scheduled to leave the platform within the next four weeks, whether it is available in Canada or the United Kingdom, and whether it was produced in India. These variables allow us to account for ownership, licensing windows, international availability, and regional content patterns. Beginning in 2024, we also use the Movie of the Night API (<https://www.movieofthenight.com/>) to measure whether titles are concurrently available through other major subscription streaming platforms, including Amazon Prime Video, Hulu, Disney+, HBO Max, and Apple TV+.

4.2 Netflix Official Weekly Top 10 Rankings

We collect Netflix’s official weekly Top 10 rankings from Tudum (<https://www.netflix.com/tudum/top10>). The U.S. rankings begin with the week of June 28 to July 4, 2021. Netflix publishes separate weekly lists for movies and series, with ranks from 1 to 10. The rankings are based on viewing during the corresponding Monday-to-Sunday week.⁵ For each ranked title, we record the title, rank, content type, week, and official synopsis. We

⁵In June 2023, Netflix revised its definition of views. Previously, titles were ranked using total viewing time. Under the new metric, views are calculated by dividing total time spent watching a series or movie by its runtime (Porter, 2023).

then match ranked titles to the weekly catalog using Netflix IDs and construct title-week Top 10 indicators.

These indicators serve two roles in the analysis. First, they are the dependent variable in the self-preferencing tests, where we ask whether Netflix Originals are more likely to receive prominent ranking visibility after conditioning on external popularity and other observables. Second, they define treatment timing in the event-study analysis of how Top 10 placement affects subsequent attention and viewing.


4.3 Wikipedia Page Views

We use Wikipedia page views as a proxy for title-level attention that is external to Netflix’s interface. Wikipedia traffic has been used as a measure of public interest in media markets. Mestyán et al. (2013) show that Wikipedia activity predicts box office performance, and Leung et al. (2025) show that movie-trailer releases coincide with sharp increases in Wikipedia page views and edits. Because Wikipedia traffic is publicly observable and not controlled by Netflix, it provides a useful measure of interest that is separate from the platform’s own viewership data.

The responsiveness of Wikipedia traffic to streaming availability is illustrated by *Community*, which was added to Netflix on April 1, 2020 and removed on April 1, 2024. Figure 1 plots weekly Wikipedia page views in the year before and after each event. When *Community* was added to Netflix, weekly Wikipedia views increased from fewer than 50,000 to more than 300,000. After it was removed, weekly views fell by roughly half. The example illustrates that Wikipedia activity responds to changes in platform availability and can capture attention around streaming events.

We identify English-language Wikipedia pages using automated searches that combine title names with terms such as “film,” “TV series,” and “Wikipedia.” We manually validate matches, especially for common titles, remakes, and non-English titles. Of roughly 22,000 unique catalog titles, approximately 18,000 are matched to a distinct Wikipedia page. For each matched title, we retrieve daily page views using the Wikimedia REST API and aggregate the counts to the weekly level. The resulting panel aligns Wikipedia attention with catalog availability and Top 10 status.

Figure 1: Wikipedia Weekly Page Views: *Community* (TV Series)



4.4 Device-Level Streaming and Linear TV Viewership Data

We also use device-level streaming and linear TV viewership records from Samba TV covering tens of millions of U.S. smart TVs between 2021 and 2024. Samba TV uses digital fingerprinting technology to identify viewed content and records the timestamp, platform, title information, and viewing duration for each session. The data include Netflix and other streaming platforms. Because of metadata restrictions, title-level matching is most complete for Netflix Originals, although aggregate platform viewing is observed more broadly.

From these records, we construct title-level and device-level engagement measures,

including total duration watched, the number of distinct viewing events, and Netflix views, defined as total duration divided by runtime. These data support three parts of the empirical analysis. First, they validate Wikipedia page views as a proxy for realized viewing. Second, they provide behavioral outcomes for the causal analysis of Top 10 placement. Third, they allow us to estimate heterogeneous device-level responsiveness to Top 10 rankings.

4.5 Summary Statistics

4.5.1 Market Shares of Streaming Services


Figure 2 reports streaming-platform market shares between September 2021 and October 2024 using total duration watched in the device-level data. Netflix accounts for 44.9% of observed viewing time, followed by YouTube at 23.6% and YouTube TV at 8.7%. Other major platforms account for substantially smaller shares, including Amazon Prime Video at 3.5%, CBS Network at 1.3%, Disney+ at 1.3%, HBO at 1.2%, Hulu at 0.9%, and Warner Brothers at 0.5%.

The viewing shares indicate a much larger Netflix advantage than subscriber counts alone would imply. FlixPatrol reports roughly 81 million U.S. Netflix subscribers, compared with 75 million for Amazon Prime, 54 million for Hulu, and 50 million for HBO.⁶ The gap between subscriber counts and viewing time highlights the value of behavioral data for measuring platform engagement.

Figure 3 plots Netflix viewing over time at the device-week level. Each device in the sample spends around 3.5 to 4 hours per week watching streaming content on any platform. Netflix viewing declines from about 2 hours per device-week in late 2021 and early 2022 to roughly 1 hour by late 2024. Netflix’s share of total viewing time is about 50% through 2021 and much of 2022, then falls to about 30% by late 2024. The decline appears to operate primarily through the extensive margin, since the share of devices watching any Netflix content falls from roughly 25–35% before 2023 to about 15% by late 2024.

⁶See <https://flixpatrol.com/streaming-services/subscribers/united-states/>.

Figure 2: Streaming Market Share by Network, September 2021–October 2024




4.5.2 Netflix Original Content Over Time

Figure 4 plots the share of Netflix Originals in the overall catalog and in the weekly Top 10 lists for movies and series. For movies, the catalog share of Originals rises from near zero in 2016 to about 50% by 2025. For series, the increase is larger, with the Original share exceeding 70% by the end of the sample.

The Top 10 lists, which begin in mid-2021, contain a higher share of Originals than the catalog, especially for series. More than 80% of series appearing in the Top 10 are Netflix Originals, compared with roughly 70% of series in the catalog. For movies, the Top 10 Original share is above the catalog share until 2023, after which it falls below the catalog share. These descriptive patterns are consistent with preferential visibility, but they are not conclusive because Originals may differ from licensed titles in investment, quality, audience demand, and release strategy. The regression analysis below tests whether the

Figure 3: Netflix Market Share Over Time



pattern remains after conditioning on observed popularity and title characteristics.


4.5.3 Validation of Wikipedia Page Views

We validate Wikipedia page views using the device-level data for September 2021 through October 2024. During this period, the device panel captures approximately 440,000 unique U.S. devices watching Netflix in a typical week. We aggregate the device-level records to title-week observations for Netflix Originals and compute three engagement measures: total duration watched, distinct viewing events, and Netflix views. Among 559,990 Netflix Original title-week observations available on the platform during this period, 286,586 are matched to the device-level data.

Figure 5 shows the relationship between Wikipedia page views and the three device-level measures. All three panels display strong positive correlations, indicating that Wikipedia traffic captures meaningful variation in realized viewing behavior.

Table 1 reports regressions of log device-level viewing measures on log Wikipedia page views for matched titles. The regressions include controls for U.S. and Canada availability and whether the title was produced in India. Wikipedia page views are positively and

Figure 4: Netflix Original Share in the Catalog and Top 10 Lists, 2016–2025



statistically significantly correlated with all device-level measures. The relationship is stronger for series, with elasticities between 0.34 and 0.57, compared with 0.29 to 0.34 for movies.


Table 1: OLS Regressions of Device-Level Measures on Wikipedia Views

	Movie			Series		
	log(Duration)	log(Watch Count)	log(Netflix View)	log(Duration)	log(Watch Count)	log(Netflix View)
log(Wikipedia Pageviews)	0.335*** (0.003)	0.290*** (0.003)	0.330*** (0.003)	0.455*** (0.003)	0.340*** (0.002)	0.565*** (0.004)
Constant	9.131 (0.049)	2.033 (0.035)	0.440 (0.047)	8.741 (0.052)	1.609 (0.039)	-0.648 (0.059)
Observations	115014	115014	114855	139252	139252	74030
R^2	0.085	0.113	0.093	0.145	0.145	0.215

Standard errors in parentheses

* $p < .1$, ** $p < .05$, *** $p < .01$

Figure 5: Wikipedia Page Views and Device-Level Viewership Metrics for Netflix Originals




We also examine where Top 10 titles rank in the weekly Wikipedia distribution. Figure 6 shows that nearly 50% of Top 10 movies and close to 70% of Top 10 series fall in the top 50 titles by weekly Wikipedia traffic within their content type. This alignment supports the use of Wikipedia page views as a scalable proxy for underlying title-level attention.

4.5.4 Top 10 and Non-Top 10 Titles

Tables 2 and 3 compare Top 10 and non-Top 10 titles using Wikipedia page views and device-level streaming data. In Table 2, Top 10 movies receive about 110,000 weekly Wikipedia views on average, compared with about 6,500 for non-Top 10 movies. For series, Top 10 titles receive about 210,000 views, compared with about 5,800 for non-Top

Figure 6: Distribution of Wikipedia Page-View Ranks Among Top 10 Titles



10 titles. The same pattern appears when attention is measured as a share of weekly Wikipedia traffic.

Table 2: Summary Statistics of Netflix Catalog Titles, Wikipedia Data

	Movie		Series	
	Non-Top 10	Top 10	Non-Top 10	Top 10
Weekly Wikipedia Pageviews (10,000)	0.6589 (2.6843)	11.0642 (18.6053)	0.5885 (1.7187)	21.4052 (37.6050)
Weekly Wikipedia Pageviews Share (%)	0.0167 (0.0628)	0.280 (0.454)	0.0151 (0.0441)	0.547 (0.926)
Netflix Original	0.410 (0.492)	0.381 (0.486)	0.706 (0.455)	0.863 (0.344)
Available (UK)	0.753 (0.431)	0.537 (0.499)	0.892 (0.311)	0.898 (0.303)
Available (Canada)	0.814 (0.389)	0.655 (0.475)	0.908 (0.289)	0.921 (0.269)
Indian Movie	0.128 (0.334)	0.0222 (0.147)	0.0219 (0.146)	0.00411 (0.0640)
N	857801	1864	475362	1568

mean coefficients; sd in parentheses

Top 10 and non-Top 10 titles also differ in composition. The share of Netflix Originals is similar across Top 10 and non-Top 10 movies, at about 40%. Among series, however, 86% of Top 10 titles are Netflix Originals, compared with 71% of non-Top 10 titles. Indian

titles are underrepresented in the Top 10, particularly for movies, where they account for 12% of all titles but only 2% of Top 10 titles.

Table 3 reports analogous statistics from the device-level data. The table reports average duration watched per device and average Netflix views per device. The sample is smaller because the device-level title match covers a shorter period and is more complete for Netflix Originals.⁷ The contrast between Top 10 and non-Top 10 titles is large in both duration and Netflix views, reinforcing that ranked titles receive substantially more behavioral engagement.

Table 3: Summary Statistics of Netflix Catalog Titles, Device-Level Streaming Data

	Movie		Series	
	Non-Top 10	Top 10	Non-Top 10	Top 10
Duration (sec per device)	0.2490 (1.1217)	34.3908 (30.1420)	0.6029 (1.9467)	59.1587 (65.5033)
Netflix Views (per device)	0.0000 (0.0002)	0.0052 (0.0042)	0.0004 (0.0024)	0.0285 (0.0619)
Available (UK)	0.9600 (0.1959)	0.9413 (0.2353)	0.9805 (0.1381)	0.9877 (0.1104)
Available (Canada)	0.9589 (0.1986)	0.9545 (0.2085)	0.9764 (0.1517)	0.9910 (0.0943)
Indian Title	0.0368 (0.1882)	0.0208 (0.1430)	0.0195 (0.1382)	0.0011 (0.0335)
N	129546	528	155619	893

mean coefficients; sd in parentheses

Finally, the Movie of the Night data show that only a small share of non-Netflix Original titles in 2024 are simultaneously available on other major subscription streaming platforms.⁸ Fewer than 1% of these titles are available on Apple TV+ or Disney+, and only 1–2% appear on Hulu, Paramount+, Peacock, or Amazon Prime Video. The overlap is somewhat higher for series than for movies. This limited overlap helps interpret Netflix catalog availability as a meaningful measure of access for most titles in the sample.

⁷The device-level title match is smaller than the catalog-Wikipedia panel because the device-level data cover a shorter period and include reliable title-level matches primarily for Netflix Originals.

⁸Almost no Netflix Original titles are available through other platforms' subscription plans.

5 Detecting Self-Preferencing

In this section, we examine whether Netflix original titles are more likely to be featured in the weekly Top 10 list than non-original titles, conditional on their popularity and other observables. The central challenge is to distinguish between promotion driven by consumer demand and differential visibility that reflects platform incentives. To address this, we compare the likelihood that Netflix Originals and non-original titles appear in the Top 10 after conditioning on observable measures of popularity and title characteristics.

5.1 Empirical Specification

We use the conditioning on observables (COO) approach defined in Reimers and Waldfogel (2023) to test whether Netflix Originals are more likely to be featured in the weekly Top 10 list than non-original titles, conditional on their popularity and other observables. Specifically, we estimate the following regression:

$$\begin{aligned} TopTen_{it} = & \beta_0 + \beta_1 X_{it}^{wiki} + \beta_2 X_{it}^{wiki} \times X_i^{original} + \beta_3 X_{it}^{wiki} \times X_{it}^{Expiring} \\ & + \sum_{c \in C} \beta_{4c} X_{it}^{wiki} \times X_{ict}^{Available} + \beta_5 X_{it}^{wiki} \times X_i^{India} + \phi_t + \epsilon_{it}, \end{aligned} \quad (1)$$

where $TopTen_{it}$ is a binary indicator for whether title i makes the Top 10 in week t , X_{it}^{wiki} is the natural log of Wikipedia page views the title receives that week, and $X_i^{original}$ indicates whether the title is a Netflix original. The interaction term $X_{it}^{wiki} \times X_i^{original}$ captures whether equally popular Netflix originals are more likely to enter the Top 10, providing a direct test of self-preferencing. We also interact page views with whether the title is expiring in the next four weeks, availability indicators for Canada and the U.K., and with an indicator for Indian titles, in order to account for streaming availability, regional availability, and cultural targeting. Week fixed effects ϕ_t absorb week-to-week variation in the overall competitiveness and visibility of titles on the platform.

5.2 Main Results

Table 4 presents the regression results. The first three columns report estimates for movies, while the last three columns correspond to series. Columns 2 and 5 incorporate an interaction term between Wikipedia page views and whether the title is set to expire within four weeks. Columns 3 and 6 further include controls for geographic availability and whether the title is an Indian production.⁹

Table 4: Self-Preferencing Regression Results

	(1) Movie	(2) Movie	(3) Movie	(4) Series	(5) Series	(6) Series
ln(Weekly Wikipedia Pageviews)	0.00154*** (0.0000268)	0.00152*** (0.0000272)	0.00174*** (0.0000296)	0.00242*** (0.0000476)	0.00237*** (0.0000484)	0.00244*** (0.0000575)
ln(Wikipedia Pageviews)XNetflix Original	0.0000536*** (0.0000163)	0.0000746*** (0.0000167)	0.000153*** (0.0000198)	0.000691*** (0.0000273)	0.000692*** (0.0000282)	0.000750*** (0.0000323)
ln(Wikipedia Pageviews)XExpiring in 4 Weeks		-0.000210*** (0.0000506)	-0.000312*** (0.0000508)		-0.000327** (0.000152)	-0.000328** (0.000152)
ln(Wikipedia Pageviews)XAvailable (UK)			-0.000302*** (0.0000233)			-0.000220*** (0.0000565)
ln(Wikipedia Pageviews)XAvailable (Canada)			-0.0000385 (0.0000245)			0.0000893 (0.0000588)
ln(Wikipedia Pageviews)XIndian Movie			-0.000352*** (0.0000252)			-0.000622*** (0.0000893)
Constant	-0.00875*** (0.000201)	-0.00866*** (0.000203)	-0.00830*** (0.000204)	-0.0175*** (0.000344)	-0.0170*** (0.000348)	-0.0170*** (0.000349)
Observations	711329	683316	683316	399718	380803	380803
R^2	0.00	0.00	0.01	0.01	0.01	0.01

Standard errors in parentheses

* $p < .1$, ** $p < .05$, *** $p < .01$

The results show that Wikipedia page views are a strong predictor of Top 10 appearance. Titles with higher Wikipedia traffic are significantly more likely to be featured in the Top 10. On average, a 10% increase in Wikipedia page views is associated with a 1.5–1.7% higher likelihood for movies and a 2.3–2.4% higher likelihood for series to enter the Top 10. This relationship remains robust after controlling for whether a title is set to expire within four weeks and becomes even stronger when accounting for international availability and whether the title is an Indian production.

More importantly, we find evidence consistent with self-preferencing: the interaction between Wikipedia page views and Netflix Original status is positive and statistically significant in both the movie and series regressions. This indicates that, conditional on

⁹Additional robustness checks are reported in Appendix B.

popularity, Netflix Originals are more likely to appear in the Top 10 than comparable non-original titles.¹⁰ Under the full specification in columns 3 and 6, a 10% increase in Wikipedia page views is associated with an additional 0.15% and 0.75% higher likelihood of appearing in the Top 10 for Netflix Original movies and series, respectively, relative to non-original titles. These magnitudes represent economically meaningful increases over the baseline effect of popularity, especially for series. This suggests that self-preferencing exists for both movies and series, but that Netflix may be promoting its own series content more aggressively, consistent with its broader strategic shift toward original serialized programming.

We also find that the interaction between Wikipedia page views and whether a title is set to expire within four weeks is negative and statistically significant. This suggests that Netflix may be systematically deprioritizing or punishing content that is about to be removed from the platform, even when those titles are attracting substantial viewer interest. Such a pattern is consistent with strategic choices in how Netflix allocates visibility to content with limited remaining licensing windows.¹¹

5.3 Heterogeneity and Dynamics


We next investigate how self-preferencing behavior evolves over time by estimating the interaction effect between Wikipedia page views and Netflix Original status separately for each half-year period. Figure 7 plots the estimated coefficients, along with 95% confidence intervals, for movies and series.

The results show that the self-preferencing effect is consistently positive and statistically significant across most years for both movies and series, but the pattern differs by content type. For series, the interaction effect remains positive and statistically significant throughout the sample period, indicating persistent preferential promotion of Netflix Original series. For movies, by contrast, the effect declines over time and disappears beginning in 2023.

¹⁰We validate Wikipedia page views as a proxy for Netflix demand using streaming device data. See Appendix A for details.

¹¹We also estimate alternative specifications where the interaction term is defined based on titles expiring within 1 to 8 weeks instead of just 4 weeks. The results, shown in Appendix B.5, remain qualitatively similar.

Figure 7: Estimated Self-Preferencing Effects Over Time: Interaction Between Wikipedia Views and Netflix Original Status (with 95% Confidence Intervals)



This change coincides with Netflix’s revision of its viewership metric in June 2023. Prior to that date, titles were ranked based on total hours viewed. Netflix then shifted to a normalized measure of views, defined as total time spent watching a series or movie divided by its runtime. Netflix explained that this new metric is anchored in engagement and provides a more comparable measure across titles of different lengths (Porter, 2023). This change may have altered the platform’s internal promotion strategy and reduced the relative advantage previously enjoyed by Netflix Original movies. The timing also coincides with the decline in Netflix’s market share documented in Figure 3, suggesting that broader competitive pressures may have contributed to a shift in promotional behavior.

Taken together, these findings indicate that self-preferencing is not constant over time. It is stronger and more persistent for series than for movies, which is consistent with the idea that original serialized content is especially valuable to Netflix as a means of retaining subscribers and reducing dependence on licensed programming.

6 Mechanisms: Why Self-Preferencing?

A central motivation for Netflix to engage in self-preferencing, by prominently recommending its own original content, stems from the risks associated with reliance on licensed,

non-original shows. Unlike Originals, which Netflix fully controls, licensed content is vulnerable to contract expiration and the growing tendency of rival streaming platforms to reclaim their own titles. When such shows leave, Netflix faces the possibility of heightened subscriber churn. For example, analysts estimated that the departure of *The Office* in 2021 alone could cost Netflix nearly \$935 million in annual subscriptions (Shuman, 2021). This illustrates the strategic importance of retaining viewers through content that cannot be withdrawn by competitors.

By investing in and preferentially promoting original programming, Netflix can reduce its exposure to these risks. Originals ensure long-term catalog stability, and promoting them through the recommendation algorithm further entrenches user habits around content that Netflix fully controls. In this sense, self-preferencing is not only about favoring in-house productions but also about safeguarding the subscription base against the volatility of licensing arrangements.

To examine the importance of this mechanism, we study whether the exit of licensed titles affects subscriber churn. If the departure of popular non-original content raises unsubscribe rates, this would provide a clear incentive for Netflix to steer attention toward Originals and thereby reduce reliance on content that may disappear from the platform.

6.1 Empirical Specification

We leverage device-level streaming data in conjunction with Netflix’s catalog and Wikipedia pageview statistics to estimate whether the exit of popular non-original content affects the likelihood of churn. Specifically, we estimate the following regression model:

$$\begin{aligned} \text{Unsubscribe}_{it} = & \sum_{k=0}^K \beta_{1k}^{\text{movies}} \ln(\text{Views}_{t-k}^{\text{movies}}) + \sum_{k=0}^K \beta_{1k}^{\text{series}} \ln(\text{Views}_{t-k}^{\text{series}}) \\ & + \beta_2 \ln(X_{it}^{\text{duration}}) + \beta_3 \ln(X_{it}^{\text{week}}) + \phi_i + \epsilon_{it}, \end{aligned} \quad (2)$$

where the dependent variable Unsubscribe_{it} is an indicator for whether device i unsubscribes in week t , defined as a cessation of Netflix viewing for at least ten consecutive weeks. The key regressors, $\text{Views}_{t-k}^{\text{movies}}$ and $\text{Views}_{t-k}^{\text{series}}$, capture cumulative Wikipedia views of movies and series exiting Netflix in week $t - k$, which serve as proxies for the

popularity of departing non-original titles. Since the device-level data do not report the specific non-original titles viewed, Wikipedia pageviews provide an external measure of audience interest in departing content.

Lagged values are included to capture delayed responses, as cancellations may occur weeks after the exit event. The specification also includes controls for cumulative watch duration ($X_{it}^{duration}$) and the number of active weeks (X_{it}^{week}) for each device. Device fixed effects (ϕ_i) absorb time-invariant heterogeneity in viewing intensity and subscription behavior.

This framework allows us to assess whether the loss of popular licensed content, particularly series, significantly increases the likelihood of unsubscribing. A strong positive relationship would support the view that self-preferencing of Originals serves as a strategic response to catalog instability.

6.2 Results

Table 5 reports the estimates of the relationship between content exits and subscriber churn. The first specification includes only the cumulative Wikipedia views of movies and series exiting Netflix in the current week, the second adds lagged exits, and the third further incorporates device-level controls such as cumulative viewing duration and the number of active weeks.

The results yield two main findings. First, the effect of movie exits on the unsubscribe rate is negative, which runs against the idea that the loss of popular films materially increases churn. The magnitude of this effect is small, however. In the full specification, a 10% increase in the cumulative Wikipedia views of exiting movies reduces the unsubscribe rate over the next three weeks by less than 0.1%, compared to an average churn rate of roughly 4% in the sample.

Second, the effect of series exits is positive and substantially larger. In the full specification, a 10% increase in the cumulative Wikipedia views of exiting series raises the unsubscribe rate over the next three weeks by about 0.4%. This result suggests that the loss of popular serialized content meaningfully contributes to subscriber churn. Relative to movies, series appear to play a much more important role in sustaining continued

Table 5: Exiting Titles on Unsubscribe Rates

	(1) Unsubscribe	(2) Unsubscribe	(3) Unsubscribe
$\ln(\text{Views}_t^{\text{movie}})$	-0.000167*** (0.0000162)	0.000186*** (0.0000179)	-0.000154*** (0.0000179)
$\ln(\text{Views}_t^{\text{series}})$	0.000812*** (0.0000172)	0.000957*** (0.0000174)	0.00111*** (0.0000175)
$\ln(\text{Views}_{t-1}^{\text{movie}})$		0.000280*** (0.0000172)	-0.0000431** (0.0000172)
$\ln(\text{Views}_{t-2}^{\text{movie}})$		0.000159*** (0.0000175)	-0.000119*** (0.0000175)
$\ln(\text{Views}_{t-3}^{\text{movie}})$		-0.000256*** (0.0000172)	-0.000566*** (0.0000172)
$\ln(\text{Views}_{t-1}^{\text{series}})$		0.000892*** (0.0000183)	0.000878*** (0.0000183)
$\ln(\text{Views}_{t-2}^{\text{series}})$		0.000913*** (0.0000178)	0.000869*** (0.0000179)
$\ln(\text{Views}_{t-3}^{\text{series}})$		0.000987*** (0.0000178)	0.000952*** (0.0000179)
$\ln(\text{Duration})$			-0.0186*** (0.0000631)
# of Weeks Active			0.000665*** (0.00000170)
Constant	0.0368*** (0.0000622)	0.0287*** (0.000189)	0.245*** (0.000789)
Observations	61611833	61611833	61611833
Device FE	Yes	Yes	Yes
# Device	1757660.00	1757660.00	1757660.00
R^2	0.12	0.12	0.13

Standard errors in parentheses

* $p < .1$, ** $p < .05$, *** $p < .01$

engagement on the platform.

The controls behave as expected. Conditional on the number of active weeks, devices with greater cumulative viewing duration are less likely to unsubscribe, while conditional on total viewing duration, devices with longer tenure on Netflix are more likely to churn. Overall, these findings highlight the strategic importance of original series for subscriber retention. They also provide a plausible mechanism for the self-preferencing patterns documented in Section 5: if the exit of popular licensed series increases churn, Netflix has a clear incentive to direct viewer attention toward Original series that remain under its control.

7 The Causal Impact of Platform Visibility

This section estimates the causal impact of appearing in Netflix’s weekly Top 10 list on subsequent title-level viewership. A central challenge is endogeneity: titles that are already gaining traction are more likely to be selected into the Top 10, so raw correlations confound the effect of visibility with underlying demand. To isolate the causal effect of Top 10 inclusion, we implement a staggered difference-in-differences event-study design. The design compares titles at the time they first enter the Top 10 with comparable titles that have not yet entered, using cohort-specific risk sets and pre-treatment matching. This framework allows us to trace short-run dynamics around first Top 10 entry while maintaining transparent identifying assumptions.

7.1 Empirical Design

The empirical analysis is conducted at the title-week level (i, t) . Define

$$TopTen_{it} = \mathbf{1}\{\text{title } i \text{ appears in the U.S. Top 10 list in week } t\}.$$

Treatment timing is staggered, as titles enter the Top 10 in different calendar weeks. For each title, let

$$g_i = \min\{t : TopTen_{it} = 1\},$$

and set $g_i = \infty$ for never-treated titles. Event time is defined as $k = t - g_i$. We estimate effects on

$$k \in \{-4, -3, -2, -1, 0, +1, +2, +3, +4\},$$

with $k = -1$ omitted as the reference period.

Following Sun and Abraham (2021), we estimate the staggered difference-in-differences event-study specification

$$Y_{it} = \alpha_i + \lambda_t + \sum_{k \in \{-4, -3, -2, 0, +1, +2, +3, +4\}} \beta_k \mathbf{1}\{t - g_i = k\} + \varepsilon_{it},$$

where, in the Wikipedia specification, the outcome is $Y_{it} = \ln(\text{weekly Wikipedia pageviews}_{it} +$

1), and, in the streaming-device-data specification, outcomes are $Y_{it}^{dur} = \ln(\text{Duration watched}_{it} + 1)$ and $Y_{it}^{view} = \ln(\text{Netflix views}_{it} + 1)$. Accordingly, coefficients are interpreted as approximate percentage changes using $100 \times (\exp(\beta_k) - 1)$. Here, α_i and λ_t denote title and week fixed effects, respectively. The coefficient β_k measures the average difference in outcome at event time k relative to the week prior to first Top 10 inclusion ($k = -1$).

The control group in week t is the not-yet-treated risk set:

$$\mathcal{C}_t = \{j : g_j > t\} \cup \{j : g_j = \infty\}.$$

Hence, controls are either future-treated titles not yet exposed by t or never-treated titles. Already-treated titles ($g_j \leq t$) are excluded from the control pool.

To align treated and control observations in the staggered setting, estimation is implemented using stacked cohort windows. For each treatment cohort g , we construct a stack on $[g - 4, g + 4]$, include treated titles with $g_i = g$, and restrict controls to titles that remain untreated through the end of the window ($g_j = \infty$ or $g_j > g + 4$). This cohort-specific risk-set construction ensures that comparisons are made against units that are still eligible controls over the full event window.

Matching is integrated into this design before estimation. Within each cohort stack, titles are matched using only pre-treatment information, separately by content type (movies versus series) and within `Netflix_Original` strata. Nearest-neighbor matching uses two pre-treatment moments: (i) outcome level at $g - 1$ and (ii) the short pre-trend slope $Y_{i,g-1} - Y_{i,g-2}$. We additionally impose exact stratification on release or missingness status, with

$$newly_released_i(g) = \mathbf{1}\{Y_{i,g-1} \text{ is missing}\},$$

and, when implemented, missingness indicators at earlier leads,

$$missing_pre\ell_i(g) = \mathbf{1}\{Y_{i,g-\ell} \text{ is missing}\}, \quad \ell = 2, 3, 4.$$

No post-treatment information is used in matching. The event-study regression is then estimated on the matched stacked sample with fixed effects, and standard errors are clustered by title and week.

Missing outcomes are never imputed as zero. In the main implementation, we require a balanced pre-period for both treated events and matched controls, so that outcomes are observed at $k = -4, -3, -2, -1$. This addresses mechanical pre-trend patterns induced by release-timing missingness. As an alternative implementation, newly released titles are retained and post-treatment dynamics ($k \geq 0$) are estimated within the newly released stratum.

Identification requires parallel trends in untreated potential outcomes between treated cohorts and their not-yet-treated controls. We assess this condition using a joint test of pre-treatment leads,

$$H_0 : \beta_{-4} = \beta_{-3} = \beta_{-2} = 0.$$

Large and systematic lead coefficients would be evidence against the identifying assumption. As a complementary check for staggered timing, we also report cohort-time ATT aggregation in the spirit of Callaway and Sant’Anna (2021).

7.2 Results

Figures 8 and 9 report the event-study estimates using Wikipedia page views and streaming-device outcomes, respectively. Figure 8 presents the dynamic effects of first Top 10 entry on weekly Wikipedia traffic for movies and series. Figure 9 presents the corresponding effects on device-level outcomes, using both log duration watched and log Netflix views.

The estimated effects rise sharply at Top 10 entry ($k = 0$), remain elevated at $k = +1$, and then attenuate over time. In the Wikipedia sample, the effect for movies is about 1.47 log points at $k = 0$ and 1.45 log points at $k = +1$, corresponding to increases of roughly 335% and 326% in pageviews. The effect remains positive through $k = +4$, where the estimate of about 0.62 log points implies an increase of roughly 86%. For series, the effects are even larger. Estimates of 1.70 and 1.95 log points at $k = 0$ and $k = +1$ imply increases of approximately 447% and 603%, and the effect remains sizable at $k = +4$, where 1.29 log points corresponds to an increase of about 263%.

The streaming-device-data results show similarly large and persistent effects, especially for series. For $\ln(\text{Duration watched} + 1)$, the estimated effect for series is around 2.57 log points at $k = 0$ and 2.02 log points at $k = +4$, implying increases of about 1,204% and

Figure 8: Wikipedia event-study coefficients (Movie vs Series), main balanced-pre sample: $Y_{it} = \log(\text{weekly pageviews}_{it} + 1)$






Figure 9: Streaming Device data event-study coefficients (Movie vs Series), main balanced-pre sample, by outcome: $\log(\text{Duration} + 1)$ and $\log(\text{Netflix Views} + 1)$



654%, respectively. For $\ln(\text{Netflix views} + 1)$, the corresponding estimates are 2.44 at $k = 0$ and 1.64 at $k = +4$, which imply increases of roughly 1,047% and 416%. Movie estimates in the streaming-device-data sample are directionally positive but less precisely estimated in the balanced-pre sample, consistent with the smaller effective sample size in those cells.

We also examine heterogeneity by title ownership. Figure 10 reports the event-study estimates separately for Netflix Originals and non-original titles in the Wikipedia sample.

Figure 10: Wikipedia event-study coefficients by title type and original status (Original vs Non-Original), main balanced-pre sample: $Y_{it} = \log(\text{weekly pageviews}_{it} + 1)$



Both Originals and non-original titles exhibit a clear jump in viewership at the time of Top 10 entry, but the effect is substantially larger for Originals. For movies, the entry effect is about 2.64 log points for Originals, compared to 1.24 log points for non-originals. For series, the corresponding estimates are about 1.97 and 0.81 log points. The same ranking persists at $k = +1$, indicating that Top 10 visibility generates larger demand responses for Netflix Originals than for comparable non-original titles.

For newly released titles, the post-only implementation also points to positive short-run effects, especially for movies across both datasets. In the Wikipedia data, coefficients around 0.92 at $k = +1$ and 0.59 at $k = +2$ imply increases of approximately 150% and 80%. In the streaming-device-data sample, newly released movie effects are similarly positive at short horizons, with estimates around 1.03 at $k = +1$ and 1.02 at $k = +2$ for duration, implying increases of about 180%.

Overall, these results indicate that Top 10 inclusion has a large and economically

meaningful causal effect on subsequent viewership. The effect is strongest at entry, remains elevated in the following weeks, and is especially pronounced for series and for Netflix Originals. These findings suggest that platform-driven visibility does not merely reflect popularity, but actively amplifies it, reinforcing the role of interface prominence as a powerful mechanism through which Netflix can shape viewer behavior.

7.3 Heterogeneity in Device-Level Top 10 Responsiveness


The device-level estimates also allow us to examine whether responsiveness to the Top 10 ranking varies across users. We draw a 10 percent random sample from the universe of 1,790,644 devices with at least eight active Netflix weeks, yielding 179,064 devices. For each sampled device, we estimate a separate viewing regression and recover the coefficient on lagged Top 10 status. For device i , the estimating equation is

$$\text{Watch}_{ijt} = \alpha_i + \rho_i \text{Watch}_{ij,t-1} + \beta_{i,\text{Top10}} \text{Top10}_{j,t-1} + \mathbf{X}'_{ijt} \Gamma_i + \varepsilon_{ijt},$$

where Watch_{ijt} is an indicator equal to one if device i watches title j in week t , and $\text{Top10}_{j,t-1}$ indicates that title j appeared in Netflix’s Top 10 list in the previous week. The coefficient $\beta_{i,\text{Top10}}$ measures the incremental association between lagged Top 10 visibility and device i ’s probability of watching title j , conditional on lagged viewing and the controls in $\mathbf{X}_{ijt}$. Because the regression is estimated separately for each device, the comparison is made within a device’s title-week viewing choices over time rather than across devices.

An important feature of the specification is that $\mathbf{X}_{ijt}$ includes proxies for prior familiarity and preferences. These controls capture prior exposure to the same title, related titles, similar genres or categories, and recent viewing histories relevant to title j . They help distinguish the incremental effect of Top 10 visibility from a device’s pre-existing propensity to watch similar content. Without these controls, the Top 10 coefficient could partially reflect persistent taste, franchise familiarity, or genre-specific demand rather than the visibility effect of the ranking itself. The implementation used here does not include title fixed effects in the per-device regressions, so these history and exposure controls are central for addressing persistent title and genre heterogeneity.

Figure 11: Distribution of Device-Level Top 10 Responsiveness



Notes: The figure plots the distribution of the estimated device-level coefficient on lagged Top 10 status for the 10 percent random device sample. The sample contains 179,064 devices with at least eight active Netflix weeks. The random seed is 42.

Figure 11 shows substantial heterogeneity in responsiveness to the Top 10 ranking. Although the average effect remains positive in the device-level analysis, there is a sizable mass of devices with coefficients close to zero. For these devices, the Top 10 list has little measurable effect on title-level viewing decisions. The distribution therefore suggests that the Top 10 row is not a uniform demand shifter across the user base.

To better understand this heterogeneity, we relate the estimated device-level Top 10 coefficient to measures of prior viewing history. The history variables are constructed from device-level Netflix viewing records, with total duration converted to hours and active weeks measured directly from the weekly device panel. Table 6 reports separate bivariate regressions of $\beta_{i,\text{Top10}}$ on total viewing duration, weeks active, and log transformations of these measures.

Across all four specifications, the estimated relationship is negative. Devices with longer viewing histories and more active weeks exhibit smaller Top 10 coefficients. The magnitudes are modest per unit but become economically more meaningful over large differences in usage histories. In the weeks-active specification, 100 additional active weeks are associated with a reduction of about 0.0013 in $\beta_{i,\text{Top10}}$. In the duration specification,

Table 6: Top 10 Coefficient and Watching History

Dependent variable: Top 10 Coefficient (β_{top10})				
	(1) Duration	(2) Weeks Active	(3) Log(Duration)	(4) Log(Weeks Active)
Coefficient	-1.52e-06*** (2.86e-08)	-1.26e-05*** (4.49e-07)	-0.000392*** (8.95e-06)	-0.000371*** (2.12e-05)
Observations	179,064	179,064	179,064	179,064
R^2	0.0303	0.0038	0.0090	0.0018

Robust HC1 standard errors in parentheses.

Each column is a separate bivariate regression with an intercept.

Duration is measured in hours. Log specifications use $\log(x + 1)$.

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

1,000 additional viewing hours are associated with a reduction of about 0.0015. The log specifications imply similar patterns. A one-unit increase in log viewing history, roughly a 170 percent increase in the underlying history measure, is associated with a decline of about 0.00039 using total duration and 0.00037 using weeks active.

These estimates are highly precise in part because the second-stage regressions use 179,064 devices. The relevant interpretation is therefore economic as well as statistical: although the marginal effect of an additional week or hour is small, differences accumulated over long usage histories imply materially lower responsiveness among veteran users. These patterns are consistent with a signal-substitution interpretation. For users with limited histories, the generic Top 10 ranking may provide useful information about what to watch. Highly engaged users may instead rely more on personalized recommendations, accumulated viewing histories, or prior knowledge of available content. These second-stage estimates are descriptive cross-sectional correlations and should not be interpreted as causal effects of viewing history on Top 10 responsiveness.

8 Discussion and Policy Implications

Our findings contribute to the policy debate over self-preferencing by showing that platform-controlled visibility can shape demand even in a subscription environment where users face zero marginal prices for consuming individual titles. In settings such as Amazon search or app store rankings, concerns about self-preferencing often center on transaction-level distortions, such as steering consumers toward the platform’s own products at the point of purchase. In the streaming context, the mechanism is different. Because Netflix earns revenue primarily through subscriber acquisition and retention rather than per-title

transactions, the incentive to favor proprietary content operates through the allocation of attention and the management of long-run churn risk. Our results suggest that this distinction is economically important. Netflix Originals are more likely to receive prominent visibility conditional on observable popularity, and that visibility increases subsequent viewing. At the same time, the evidence does not imply that every form of platform curation is harmful. It shows that ranking and recommendation interfaces are economically consequential, and that policy discussions of platform neutrality should account for demand effects in subscription media markets as well as transaction-based markets.

9 Conclusion

This paper studies self-preferencing and ranking effects in subscription video streaming. Using a hand-constructed panel of the Netflix U.S. catalog, official Top 10 rankings, Wikipedia page views, and device-level streaming records, we examine whether Netflix allocates prominent visibility differently across its own and non-original titles and whether this visibility changes subsequent demand. The evidence is consistent with Netflix self-preferencing its own content, especially series. Conditional on external popularity and other observable controls, Netflix Originals are more likely to appear in the Top 10 than comparable non-original titles. We also show that Top 10 placement causally increases later viewership, which implies that rankings are not only summaries of contemporaneous popularity. They are also demand shifters that can amplify attention after a title receives platform-controlled visibility.

The device-level evidence adds a second contribution. Top 10 responsiveness varies substantially across users. Many devices display little measurable response to the generic ranking, while others respond strongly. Moreover, responsiveness is lower among devices with longer and richer Netflix viewing histories. This pattern is descriptive, but it is consistent with a broader economic interpretation in which public popularity signals matter more for users with limited histories, while personalized recommendations and accumulated knowledge of the catalog substitute for generic rankings among more experienced users. The result highlights a margin that is central to modern platform competition: the interaction between public interface signals and user-specific recommendation systems.

Several limitations should be kept in mind. The empirical setting is observational, even though the analysis uses external demand proxies, matched event-study designs, and device-level controls to address the main sources of confounding. We do not observe Netflix’s internal recommendation algorithm, homepage placement outside the Top 10 list, or the full set of experiments and personalization rules that determine what a user sees. The evidence is also specific to Netflix, the U.S. catalog, and the sample period we observe. Streaming platforms differ in business models, content ownership, interface design, and recommendation technology, so the magnitudes estimated here may not generalize directly to other platforms, countries, or time periods.

These limitations point to several directions for future work. Comparative studies across streaming platforms could show whether self-preferencing is a general feature of subscription media markets or a Netflix-specific pattern. Future research could also estimate the welfare effects of platform steering by combining demand responses with substitution across titles, subscriber retention, and content acquisition costs. Another important question is how personalized recommendations interact with public rankings, especially as platforms accumulate richer user histories and deploy more sophisticated algorithms. Finally, the long-run consequences for creators and studios remain open. If platform-controlled visibility systematically affects which content succeeds, then ranking transparency, non-discrimination rules, and auditing requirements may shape not only consumer choice but also investment incentives in content markets.

References

- AGUIAR, L., AND B. MARTENS (2016): “Digital Music Consumption on the Internet: Evidence from Clickstream Data,” *Information Economics and Policy*, 34, 27–43, 10.1016/j.infoecopol.2015.01.003.
- AGUIAR, L., AND J. WALDFOGEL (2018): “Netflix: Global Hegemon or Facilitator of Frictionless Digital Trade?” *Journal of Cultural Economics*, 42, 419–445, 10.1007/s10824-017-9315-z.
- (2021): “Platforms, Power, and Promotion: Evidence from Spotify Playlists,” *The Journal of Industrial Economics*, 69, 653–691, <https://doi.org/10.1111/joie.12263>.
- AGUIAR, L., J. WALDFOGEL, AND S. WALDFOGEL (2021): “Playlisting favorites: Measuring platform bias in the music industry,” *International Journal of Industrial Organization*, 78, 102765.
- ARMSTRONG, M., AND J. ZHOU (2011): “PAYING FOR PROMINENCE,” *The Economic Journal*, 121, F368–F395.
- BOURREAU, M., AND G. GAUDIN (2022): “Streaming platform and strategic recommendation bias,” *Journal of Economics & Management Strategy*, 31, 25–47.
- CALLAWAY, B., AND P. H. C. SANT’ANNA (2021): “Difference-in-Differences with Multiple Time Periods,” *Journal of Econometrics*, 225, 200–230.
- CURE, M., M. HUNOLD, R. KESLER, U. LAITENBERGER, AND T. LARRIEU (2022): “Vertical Integration of Platforms and Product Prominence,” *Quantitative Marketing and Economics*, 20, 353–395, 10.1007/s11129-022-09255-4.
- DANAHER, B., AND S. DOOSTI (2025): “Netflix and Chilling Effect? The Impact of Streaming Video Licenses on Paid Digital Downloads,” Working paper.
- DE CORNIÉRE, A., AND G. TAYLOR (2019): “A model of biased intermediation,” *The RAND Journal of Economics*, 50, 854–882.

- FARRONATO, C., A. FRADKIN, AND A. MACKAY (2023): “Self-Preferencing at Amazon: Evidence from Search Rankings,” *AEA Papers and Proceedings*, 113, 239–43, [10.1257/pandp.20231068](https://doi.org/10.1257/pandp.20231068).
- FRICK, S. J., D. FLETCHER, AND A. C. SMITH (2023): “Pirate and Chill: The Effect of Netflix on Illegal Streaming,” *Journal of Economic Behavior & Organization*, 209, 334–347, [10.1016/j.jebo.2023.03.013](https://doi.org/10.1016/j.jebo.2023.03.013).
- GUTIÉRREZ, G. (2022): “The Welfare Consequences of Regulating Amazon,” July, http://germangutierrezg.com/Gutierrez2021_AMZ_welfare.pdf, Working paper, University of Washington, Foster School of Business.
- HAGIU, A., AND B. JULLIEN (2014): “Search diversion and platform competition,” *International Journal of Industrial Organization*, 33, 48–60.
- HAGIU, A., T.-H. TEH, AND J. WRIGHT (2022): “Should platforms be allowed to sell on their own marketplaces?” *The RAND Journal of Economics*, 53, 297–327, <https://doi.org/10.1111/1756-2171.12408>.
- HUNOLD, M., R. KESLER, AND U. LAITENBERGER (2020): “Rankings of Online Travel Agents, Channel Pricing, and Consumer Protection,” *Marketing Science*, 39, 92–116, [10.1287/mksc.2019.1167](https://doi.org/10.1287/mksc.2019.1167).
- JÜRGENSMEIER, L., J. BISCHOFF, AND B. SKIERA (2024): “Opportunities for Self-Preferencing in International Online Marketplaces,” *International Marketing Review*, 41, 1118–1132, [10.1108/IMR-03-2024-0094](https://doi.org/10.1108/IMR-03-2024-0094), Open Access under CC BY 4.0.
- LAM, H. T. (2021): “Platform Search Design and Market Power,” https://www.anderson.ucla.edu/sites/default/files/document/2021-09/Lam_JMP_Platform_Search_Design.pdf, Working paper.
- LEE, K. H., AND L. MUSOLFF (2023): “Entry Into Two-Sided Markets Shaped By Platform-Guided Search,” https://lmusolff.com/papers/Entry_and_Platform_Guided_Search.pdf, Working paper.

- LEUNG, T. C., S. QI, AND K. STRUMPF (2025): “An Empirical Decomposition of Information vs Sentiment Shifts in Advertising,” Working paper.
- McKENZIE, J., P. CROSBY, J. COX, AND A. COLLINS (2019): “Experimental evidence on demand for “on-demand” entertainment,” *Journal of Economic Behavior & Organization*, 161, 98–113, <https://doi.org/10.1016/j.jebo.2019.03.017>.
- MESTYÁN, M., T. YASSERI, AND J. KERTÉSZ (2013): “Early Prediction of Movie Box Office Success Based on Wikipedia Activity Big Data,” *PLOS ONE*, 8, 1–8, [10.1371/journal.pone.0071226](https://doi.org/10.1371/journal.pone.0071226).
- PORTER, R. (2023): “Netflix’s New View Metric Is Easier to Grasp, but It Still Doesn’t Say a Lot,” *The Hollywood Reporter*, <https://www.hollywoodreporter.com/tv/tv-news/netflix-viewership-data-metric-explained-1235521282/>, Accessed: 2025-05-09.
- REIMERS, I., AND J. WALDFOGEL (2023): “A Framework for Detection, Measurement and Welfare Analysis of Platform Bias,” Working Paper.
- SHUMAN, T. (2021): “Loss of {The Office} Could Cost Netflix \$935 Million in Annual Subscriptions,” May 13, <https://www.inmyarea.com/research/netflix-losing-subscribers-when-the-office-leaves>, InMyArea Research.
- SUN, L., AND S. ABRAHAM (2021): “Estimating Dynamic Treatment Effects in Event Studies with Heterogeneous Treatment Effects,” *Journal of Econometrics*, 225, 175–199.
- TENG, X. (2023): “Self-preferencing, Quality Provision, and Welfare in Mobile Application Markets,” <https://drive.google.com/file/d/1El3Staxp8-Z8XKgt1R1Su14ia4YfQWv7/view>, Working Paper.
- WALDFOGEL, J. (2024): “Amazon Self-preferencing in the Shadow of the Digital Markets Act,” working paper, National Bureau of Economic Research, [10.3386/w32299](https://doi.org/10.3386/w32299).
- YUAN, Y. (2025): “Exclusive Contracts in the Video Streaming Market,” https://yihao-yuan.github.io/MyWebsite/streaming_exclusivity.pdf, Working paper, The Wharton School, University of Pennsylvania.

Appendix A Validation of Wikipedia Page Views Using Streaming Device Data

To validate Wikipedia page views as a proxy for demand, we estimate a Top 10 specification on the subset of Netflix Original titles that can be matched in both the Wikipedia and streaming-device datasets. The estimating equation is

$$\begin{aligned} TopTen_{it} = & \beta_0 + \beta_1 X_{it}^{wiki} + \beta_3 X_{it}^{wiki} \times X_{it}^{Expiring} \\ & + \sum_{c \in C} \beta_{4c} X_{it}^{wiki} \times X_{ict}^{Available} + \beta_5 X_{it}^{wiki} \times X_i^{India} + \phi_t + \epsilon_{it}. \end{aligned} \quad (A1)$$

Because the streaming-device sample contains reliably matched title-level observations primarily for Netflix Originals, we exclude the interaction term $\beta_2 X_{it}^{wiki} \times X_i^{original}$ from this validation exercise. We then estimate the same specification twice, once using Wikipedia page views and once using streaming-device views. To ensure comparability, both regressions are run on the same intersection sample of title-weeks, so the number of observations is identical across data sources within content type. In the streaming-device regressions, views are measured using view counts. Results based on duration watched are similar.

Table A1: Validation of Wikipedia Proxy Using Streaming Device Data

	Movie (Wikipedia)	Movie (Streaming)	Series (Wikipedia)	Series (Streaming)
ln(views)	0.0036*** (0.0002)	0.0035*** (0.0003)	0.0051*** (0.0003)	0.0048*** (0.0004)
ln(views)XExpiring in 4 Weeks	0.0009* (0.0005)	-0.0033*** (0.0010)	-0.0008 (0.0009)	-0.0013 (0.0017)
ln(views)XAvailable (UK)	-0.0002* (0.0001)	0.0007*** (0.0003)	0.0005* (0.0002)	-0.0005 (0.0004)
ln(views)XAvailable (Canada)	0.0003*** (0.0001)	0.0007** (0.0003)	0.0007*** (0.0002)	0.0021*** (0.0004)
ln(views)XIndian Movie	-0.0004*** (0.0001)	-0.0005 (0.0004)	-0.0009*** (0.0003)	-0.0016** (0.0008)
Observations	114855	114855	74030	74030
R ²	0.015	0.032	0.026	0.039

Standard errors in parentheses
* $p < .1$, ** $p < .05$, *** $p < .01$

Table A1 shows that the log-views coefficient is nearly identical when using Wikipedia page views and streaming-device views. The estimates are 0.0036 and 0.0035 for movies, and 0.0051 and 0.0048 for series. Economically, a 1% increase in views implies nearly the same increase in the probability of Top 10 inclusion across the two data sources. The

streaming-device regressions have higher R^2 values, 0.032 versus 0.015 for movies and 0.039 versus 0.026 for series, indicating that Wikipedia page views are useful but noisier than direct viewing measures.

Appendix B Robustness Checks

This appendix reports robustness checks for the self-preferencing results in Section 5. We vary the popularity measure, restrict the estimation sample to high-attention titles, control for availability on other streaming platforms, and vary the expiration-window definition. Across these exercises, the positive relationship between title popularity and Top 10 inclusion remains robust. Netflix Originals also remain more likely to appear in the Top 10 than non-original titles with similar observable popularity, consistent with the main results.

Appendix B.1 Actual Views

As a first robustness check, we replace log Wikipedia page views with raw page view counts. The results, reported in Table B1, are qualitatively similar to those in Table 4, though the estimated self-preferencing effects are larger. In this specification, among titles with the same number of Wikipedia page views, Netflix Original movies and series are roughly twice as likely to appear in the Top 10.

Table B1: Self-Preferencing Regression Results, Actual Views

	(1) Movie	(2) Movie	(3) Movie	(4) Series	(5) Series	(6) Series
Weekly Wikipedia Pageviews (10,000)	0.00150*** (0.0000256)	0.00172*** (0.0000282)	0.00312*** (0.0000459)	0.00342*** (0.0000786)	0.00344*** (0.0000808)	0.00670*** (0.000166)
Wikipedia PageviewsXNetflix Original	0.00419*** (0.0000401)	0.00410*** (0.0000420)	0.00483*** (0.0000479)	0.00499*** (0.0000831)	0.00490*** (0.0000854)	0.00627*** (0.000101)
Wikipedia PageviewsXExpiring in 4 Weeks		-0.00159*** (0.0000705)	-0.000976*** (0.0000739)		-0.00344*** (0.000521)	-0.00276*** (0.000521)
Wikipedia PageviewsXAvailable (UK)			-0.00197*** (0.0000690)			-0.00330*** (0.000159)
Wikipedia PageviewsXAvailable (Canada)			-0.000167** (0.0000708)			-0.00135*** (0.000189)
Wikipedia PageviewsXIndian Movie			-0.00132*** (0.0000749)			-0.000687 (0.000484)
Constant	0.000421*** (0.0000598)	0.000324*** (0.0000609)	0.000263*** (0.0000609)	-0.000810*** (0.0000897)	-0.000757*** (0.0000922)	-0.000893*** (0.0000924)
Observations	711329	683316	683316	399718	380803	380803
R^2	0.05	0.05	0.05	0.17	0.17	0.17

Standard errors in parentheses

* $p < .1$, ** $p < .05$, *** $p < .01$

Appendix B.2 Top 300 Titles

In the second robustness check, we restrict the sample to the top 300 titles in each week based on Wikipedia page views. The results are reported in Table B2. This restriction reduces the sample from approximately 700,000 observations for movies and 400,000 for series in the full sample to about 55,000 to 58,000 observations for each content type. The estimated coefficients remain close to those in Table B1. This result indicates that the main findings are not driven by noise in the lower tail of the popularity distribution.

Table B2: Self-Preferencing Regression Results, Top 300 Titles

	(1) Movie	(2) Movie	(3) Movie	(4) Series	(5) Series	(6) Series
Weekly Wikipedia Pageviews (10,000)	0.00102*** (0.0000851)	0.00121*** (0.0000946)	0.00258*** (0.000151)	0.00340*** (0.000207)	0.00340*** (0.000212)	0.00693*** (0.000431)
Wikipedia PageviewsXNetflix Original	0.00435*** (0.000125)	0.00431*** (0.000132)	0.00499*** (0.000150)	0.00504*** (0.000211)	0.00497*** (0.000216)	0.00630*** (0.000249)
Wikipedia PageviewsXExpiring in 4 Weeks		-0.00122*** (0.000221)	-0.000660*** (0.000232)		-0.00364*** (0.00127)	-0.00294** (0.00127)
Wikipedia PageviewsXAvailable (UK)			-0.00190*** (0.000221)			-0.00340*** (0.000389)
Wikipedia PageviewsXAvailable (Canada)			-0.000127 (0.000226)			-0.00145*** (0.000465)
Wikipedia PageviewsXIndian Movie			-0.00135*** (0.000238)			-0.000250 (0.00121)
Constant	0.0114*** (0.000708)	0.0109*** (0.000730)	0.0103*** (0.000731)	-0.000200 (0.000633)	0.000122 (0.000646)	-0.000964 (0.000657)
Observations	57380	54988	54988	57399	55004	55004
R^2	0.05	0.05	0.05	0.18	0.18	0.18

Standard errors in parentheses

* $p < .1$, ** $p < .05$, *** $p < .01$

Appendix B.3 View Shares

In a third robustness check, we replace page view levels with Wikipedia page view share, defined as the proportion of total Wikipedia traffic accounted for by each title in a given week. This measure absorbs week-to-week variation in overall Wikipedia usage. The results are presented in Table B3. Although the coefficient scale changes, the qualitative pattern is similar to the baseline.

Netflix Original movies and series remain significantly more likely to appear in the Top 10 than non-original titles with the same Wikipedia share. Titles with higher Wikipedia shares are more likely to enter the Top 10, expiring titles are less likely to appear, titles available in the United Kingdom and Canada are favored, and Indian movies receive less

Table B3: Self-Preferencing Regression Results, View Shares

	(1) Movie	(2) Movie	(3) Movie	(4) Series	(5) Series	(6) Series
Weekly Wikipedia Pageviews Share (%)	0.0718*** (0.00111)	0.0778*** (0.00119)	0.141*** (0.00194)	0.137*** (0.00309)	0.139*** (0.00319)	0.281*** (0.00645)
Wikipedia Pageviews ShareXNetflix Original	0.169*** (0.00167)	0.169*** (0.00174)	0.200*** (0.00200)	0.204*** (0.00327)	0.202*** (0.00337)	0.262*** (0.00398)
Wikipedia Pageviews ShareXExpiring in 4 Weeks		-0.0678*** (0.00360)	-0.0461*** (0.00371)		-0.138*** (0.0202)	-0.114*** (0.0202)
Wikipedia Pageviews ShareXAvailable (UK)			-0.0904*** (0.00278)			-0.145*** (0.00633)
Wikipedia Pageviews ShareXAvailable (Canada)			-0.00356 (0.00282)			-0.0585*** (0.00740)
Wikipedia Pageviews ShareXIndian Movie			-0.0656*** (0.00300)			-0.0321* (0.0193)
Constant	0.000179*** (0.0000599)	0.000111* (0.0000610)	0.0000351 (0.0000609)	-0.000975*** (0.0000895)	-0.000930*** (0.0000919)	-0.00108*** (0.0000921)
Observations	711329	683316	683316	399718	380803	380803
R^2	0.05	0.05	0.06	0.18	0.18	0.18

Standard errors in parentheses

* $p < .1$, ** $p < .05$, *** $p < .01$

visibility relative to their Wikipedia popularity.

Appendix B.4 Availability on Other Streaming Platforms

The final self-preferencing robustness check controls for whether a title is available on other streaming platforms besides Netflix. This analysis is limited to 2024, the period during which we can reliably observe platform availability using third-party APIs. Results are shown in Table B4. The first two columns report results for movies, and the last two report results for series.

For each content type, the first column replicates the baseline regression from Table B1 using only 2024 observations. The second column adds controls for availability on other platforms. The additional controls have little effect on the estimates. The key finding remains unchanged: Netflix Originals are more likely to appear in the Top 10 than non-original titles with similar observed popularity.

Appendix B.5 Expiration Window Definition

The main specification defines the expiration indicator as equal to one if a title is scheduled to leave Netflix within four weeks. To assess the robustness of this choice, we re-estimate the same regression while varying the expiration window from 1 week to 8 weeks. For

Table B4: Self-Preferencing Regression Results with Other-Platform Availability

	(1) Movie	(2) Movie	(3) Series	(4) Series
Weekly Wikipedia Pageviews (10,000)	0.00790*** (0.000115)	0.00808*** (0.000120)	0.0132*** (0.000269)	0.0125*** (0.000287)
Wikipedia PageviewsXNetflix Original	0.00274*** (0.000107)	0.00271*** (0.000107)	0.00607*** (0.000179)	0.00713*** (0.000206)
Wikipedia PageviewsXAvailable (UK)	-0.00172*** (0.000186)	-0.00197*** (0.000189)	-0.00327*** (0.000378)	-0.00412*** (0.000399)
Wikipedia PageviewsXAvailable (Canada)	-0.00411*** (0.000175)	-0.00400*** (0.000175)	-0.00814*** (0.000398)	-0.00777*** (0.000408)
Wikipedia PageviewsXIndian Movie	0.000547*** (0.000146)	0.000519*** (0.000147)	-0.00726*** (0.000883)	-0.00728*** (0.000882)
Wikipedia PageviewsXExpiring in 4 Weeks	-0.00491*** (0.000387)	-0.00476*** (0.000387)	-0.00443** (0.00223)	-0.00443** (0.00223)
Wikipedia PageviewsXApple Subscription		0.000612 (0.000489)		-0.00719 (0.00776)
Wikipedia PageviewsXDisney Subscription		-0.00765 (0.0157)		-0.00450 (0.0236)
Wikipedia PageviewsXHulu Subscription		-0.00311*** (0.00110)		-0.00220*** (0.000516)
Wikipedia PageviewsXParamount Subscription		-0.00505*** (0.000545)		0.00255*** (0.000246)
Wikipedia PageviewsXPeacock Subscription		-0.00435*** (0.000786)		0.00476*** (0.000583)
Wikipedia PageviewsXPrime Subscription		0.00196*** (0.000521)		0.00703*** (0.000918)
Constant	-0.000469*** (0.000117)	-0.000455*** (0.000117)	-0.00121*** (0.000155)	-0.00116*** (0.000155)
Observations	192265	192265	120834	120834
R^2	0.06	0.06	0.16	0.16

Standard errors in parentheses

* $p < .1$, ** $p < .05$, *** $p < .01$

each value of k , the expiration indicator equals one if a title is scheduled to leave the platform within k weeks. Figure B1 plots the resulting interaction coefficients.

The results are stable across specifications. Expiring titles are less likely to be featured in the Top 10 regardless of whether expiration is defined as occurring within 1, 2, or up to 8 weeks. The magnitude is stable and statistically significant for both movies and series across most expiration windows. The main exception is the movie estimate at the 8-week window, where the coefficient becomes statistically insignificant, consistent with weaker expiration effects at longer horizons.

Figure B1: Effect of Expiring by Weeks and Content Type

